


## 7 Language Models

- Sequence Modeling
  - Recurrent Neural Networks (RNN)
  - Long Short-Term Memory (LSTM)
  - Example: LSTMs in Machine Translation
  - Embeddings From Language Models (ELMo)
- Transformers
  - Elements of Transformers
  - Example: Step-By-Step Calculation
- Bidirectional Encoder Representations from Transformers (BERT)
  - Training BERT
  - Fine-Tuning BERT
- BERT-alike Models
  - Sentence-BERT
  - BERTopic
  - Decoding-enhanced BERT with Disentangled Attention (DeBERTa)
- Large Language Models


# Sequence modeling

- e.g., sentiment classification for movie reviews
- word2vec still BOW
- sequence (order) (in most implementations) neglected


one to one




one to many




many to one



many to many



many to many




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>


# Recurrent neural network (RNN)

- idea: pass information from previous inputs into current calculation
- below: just a snippet of one stream of information, not a full NN!


NN representation



unfolded representation




# RNN in comparison to FF NN




- traditional RNN structure; also known as
  - simple recurrent networks
  - Elman network
  - Jordan network
- more complex architectures possible

# RNN architectures




# RNN architectures



- many to many, seq2seq
- advantage: short-term **memory**
- disadvantage:

# RNN architectures




- many to many, seq2seq
- advantage: short-term **memory**
- disadvantage: **short-term** memory

# One recurrent unit

- simplify notation:  $h_i^{(t)}, i = 1, \dots, n_1 \rightarrow h_t$ 
  - analogously we use  $x_t$

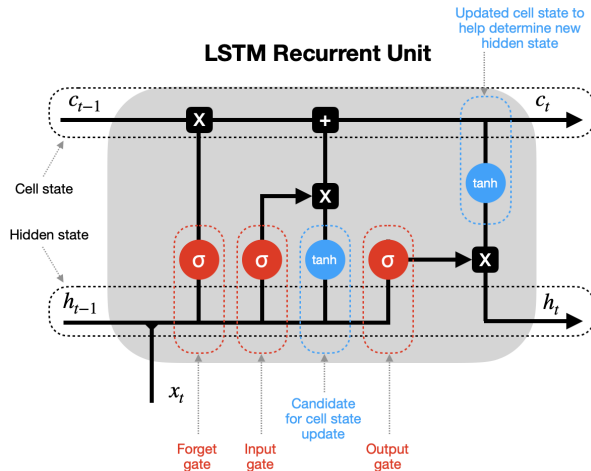
## Standard Recurrent Unit



<https://towardsdatascience.com/lstm-recurrent-neural-networks-how-to-teach-a-network-to-remember-the-past-55e54c2ff22e>



# Long short-term memory (LSTM)




<https://towardsdatascience.com/lstm-recurrent-neural-networks-how-to-teach-a-network-to-remember-the-past-55e54c2ff22e>

# LSTM in real world

- Siri
- Apple's auto-completion
- AlphaGo (software for the board game Go)
- GoogleTranslate was based on LSTMs, besides other systems
  
- no benefits in one to one scenarios
- (were) popular in seq2seq scenarios

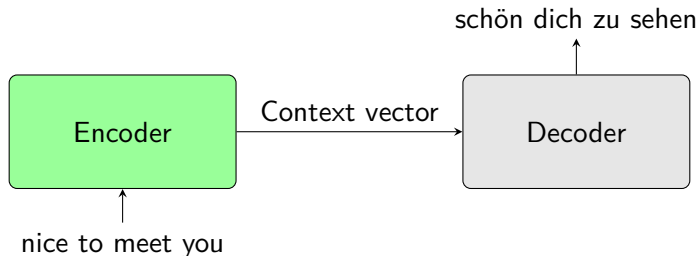
# Encoder-decoder

- how to solve machine translation (MT) problem
- idea of encoder-decoder models:
  - ① encode on input sequence
  - ② decode to output sequence
- use, e.g., 2 LSTMs one after another
- concept used for (all) sota models: BERT, GPT-3, ...

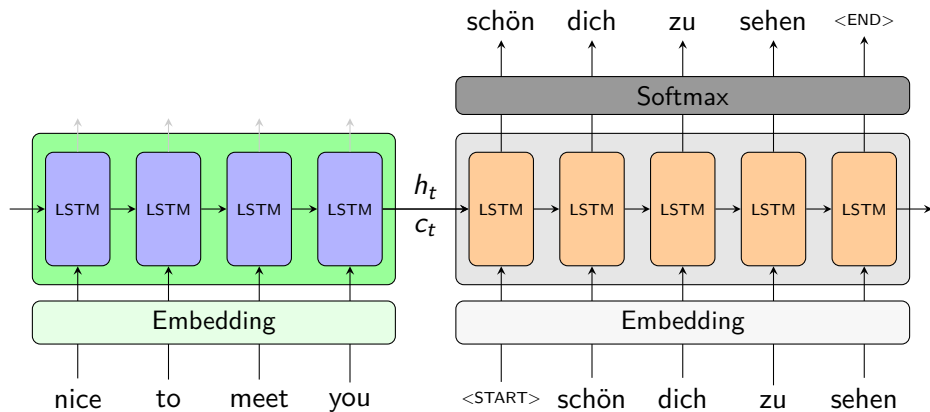


# Encoder-decoder in MT

- input: nice to meet you
- task: translation to German
- (target) output: schön dich zu sehen
- encoder: collect information about input in context vector
- context vector: final state of encoder, encapsulates meaning of input
- decoder: predict output token by token using context vector



# Example: encoder-decoder LSTM in MT

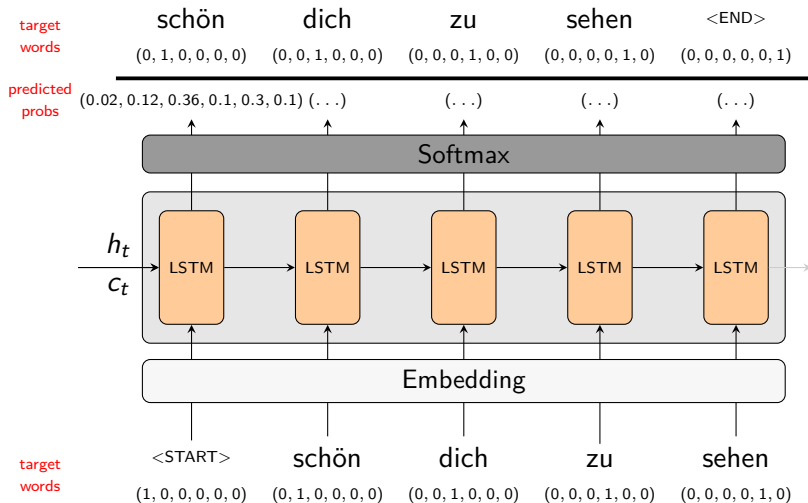


cf. <https://medium.com/analytics-vidhya/encoder-decoder-seq2seq-models-clearly-explained-c34186fbf49b>

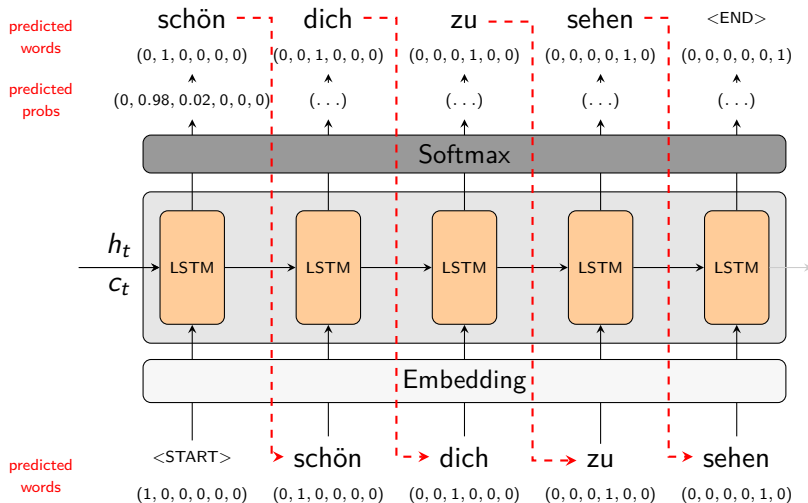
# Training and testing phases

- vectorizing the tokens: one-hot-encoding → embeddings
- <START> and <END> also get one-hot-encoding & embeddings
- training & testing of encoder identical and straightforward
- instead, for decoder two different non-trivial approaches:
  - teacher forcing while training (cf. next slide):  
use true labels as inputs, not the predicted output sequence
  - during the testing phase the predicted token are used as input for next token — as one-hot-encoded vector passed to the embedding layer

# Decoder in the training phase

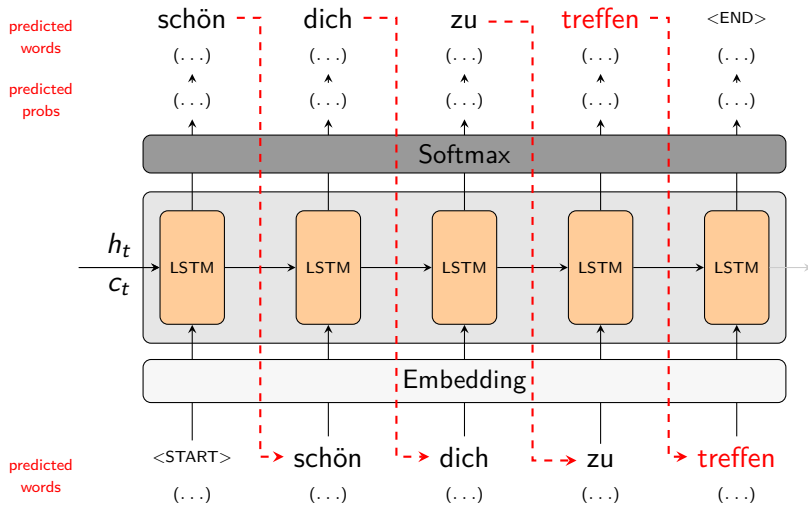


# Decoder in the testing phase





# Wrong predictions in testing phase



# Embeddings from language models (ELMo)

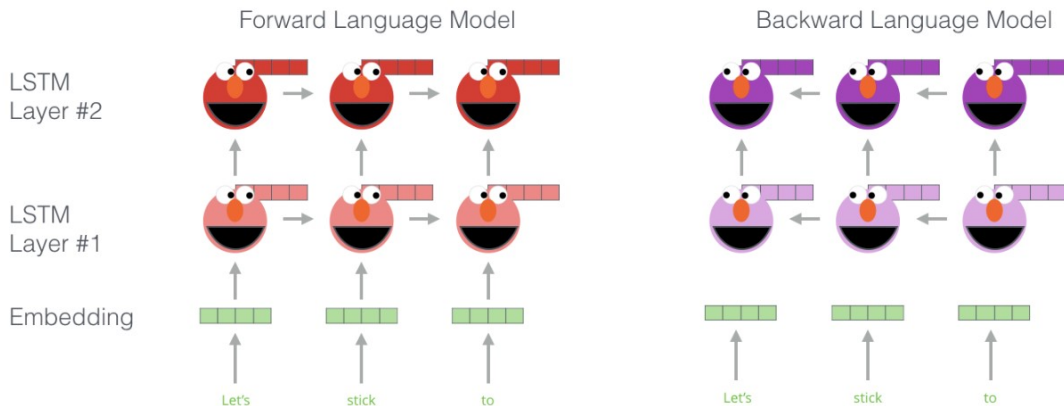
Let's stick to statistics.

The dog gets the stick.

- learn different embeddings for words in different contexts
- word2vec, fastText, GloVe would all result in the same embeddings for both appearances
- ELMo: based on LSTM layers
- idea: encode each word including its context
- first model using contextual embeddings
- bidirectional LSTM with LM objective (predict word)

## ELMo — step 1

## Embedding of “stick” in “Let’s stick to” - Step #1



<https://jalammr.github.io/illustrated-bert/>

# ELMo — step 2

## Embedding of “stick” in “Let’s stick to” - Step #2

1- Concatenate hidden layers



2- Multiply each vector by a weight based on the task




3- Sum the (now weighted) vectors




ELMo embedding of “stick” for this task in this context

Forward Language Model



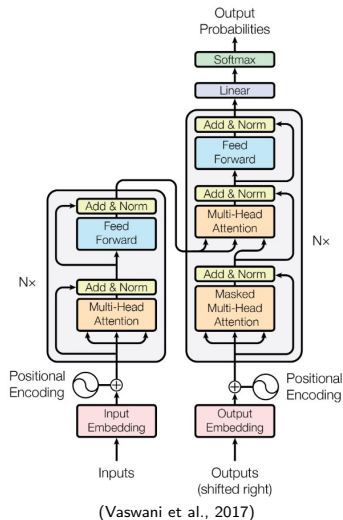
Backward Language Model



<https://jalammr.github.io/illustrated-bert/>

# Transformer idea

- RNN cannot be computed in parallel
- transformer enables parallel computation by using
  - attention
  - positional encoding
  - FF layers
- replacement of LSTMs
- better long-term memory




# Positional encoding

$$\text{PE}(x \mid \text{pos}) = \begin{cases} \sin(\text{pos}/10000^{x/d}) & \text{if } x \text{ is even} \\ \cos(\text{pos}/10000^{(x-1)/d}) & \text{if } x \text{ is odd} \end{cases}$$

- enables parallel calculation
- is simply added to the word embedding
- pos is the token position  $(0, \dots)$
- $x$  is the embedding dimension  $(0, \dots, d)$

# Positional encoding (visualization)

$$\text{PE}(x \mid \text{pos}) = \begin{cases} \sin(\text{pos}/10000^{x/d}) & \text{if } x \text{ is even} \\ \cos(\text{pos}/10000^{(x-1)/d}) & \text{if } x \text{ is odd} \end{cases}$$



<https://jalammar.github.io/illustrated-transformer/> — here  $d = 64$

# Attention is all you need<sup>1</sup>

## Self-attention

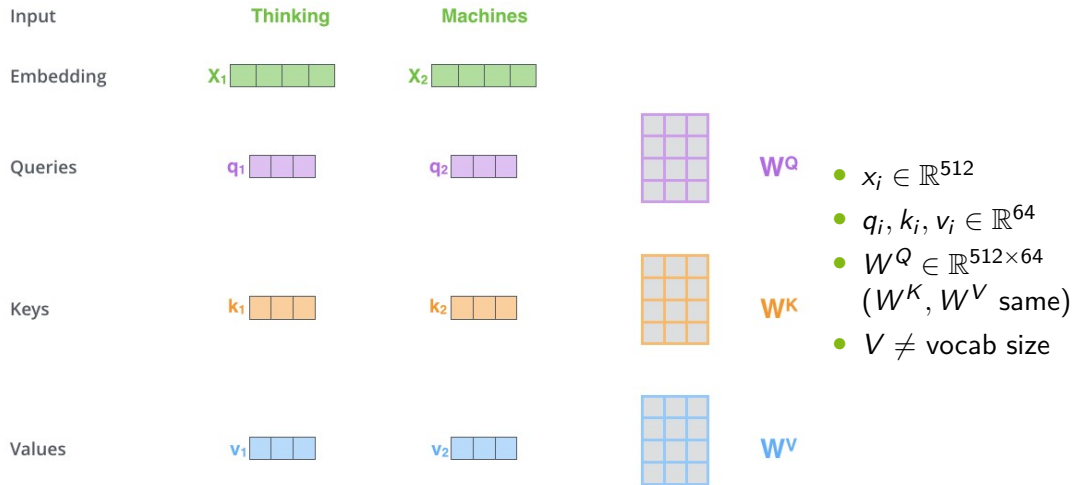
The trophy does not fit into the suitcase because it is too big/small. (cf. slide 29)

- question: what does "it" refers to in the sentence above
- human: if big → trophy, if small → suitcase
- not that easy for LM
- self-attention allows to look at other tokens' embeddings
- similar approach as in RNNs

<sup>1</sup>Attention Is All You Need (Vaswani et al., 2017) <https://doi.org/10.48550/arXiv.1706.03762>

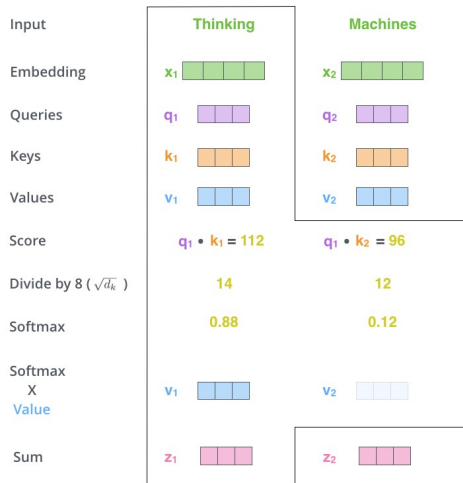


# Self-attention elements



<https://jalammar.github.io/illustrated-transformer/>

# Self-attention calculation



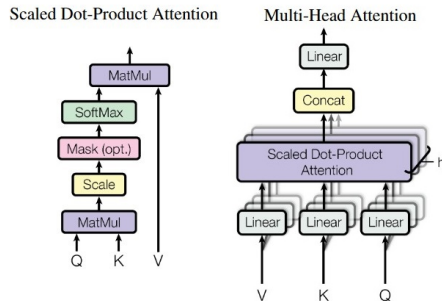
- 1 get  $q_1, k_1, v_1$  by multiplying  $x_1$  with  $W^Q, W^K, W^V$
- 2 dot-product  $q_1$  and  $k_i$  for all  $i$  in the same sequence
- 3 divide by 8
- 4 calculate softmax
- 5 weighted sum of  $v_i$
- 6  $z_1$  resulting attention vector

<https://jalammar.github.io/illustrated-transformer/>

# Self-attention matrix representation

- $Q = XW^Q$
- $K = XW^K$
- $V = XW^V$
  
- scaled dot-product attention
- $Z = \text{softmax}(QK^T / \sqrt{64})V$
- 64 because it is the respective dimension  $d_k$   
⇒ to prevent from extremely small gradients for the softmax function
  
- in general:
  - $d \times d_k$  dimension of  $W^Q$  and  $W^K$
  - $d \times d_v$  dimension of  $W^V$
  - $d_k = d_v = d/h = 64$  often selected ( $d = 512$ )
  - $h = 8$  number of (multi-head) attention layers (cf. following slides)

# (Simple) self-attention and multi-head attention (MHA)



(Vaswani et al., 2017)

- idea: several (typically  $h = 8$ ) parallel self-attention layers
- each layer returns an attention matrix  $Z_i, i = 1, \dots, h$  (cf. previous slide)
- $Z = (Z_1, \dots, Z_h)W^0$
- $W^0 \in \mathbb{R}^{hd_v \times d} (= \mathbb{R}^{512 \times 512})$

# MHA in one figure


1) This is our input sentence\*

Thinking  
Machines

2) We embed each word\*



3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices





4) Calculate attention using the resulting  $Q/K/V$  matrices



5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer




\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one




...

...

...




$W^O$



<https://jalammar.github.io/illustrated-transformer/>

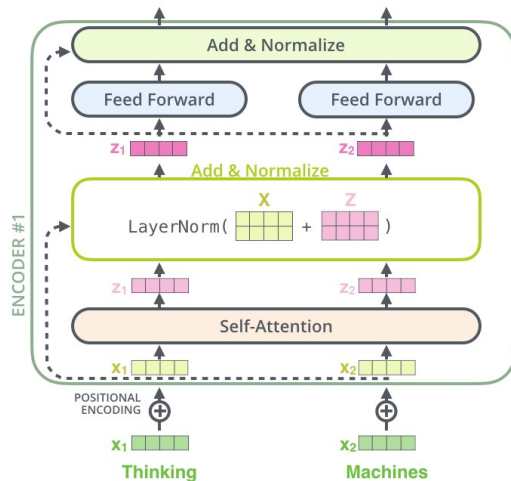
## Now we're finally encoding

- on the right: first encoder layer
- all other layers get  $r$  vectors
- $\text{FFNN}(z) = \max(0, zW_1 + b_1)W_2 + b_2$   
(two linear transformations + ReLU)
- typically:
  - $x_i \in \mathbb{R}^{512}$
  - $z_i \in \mathbb{R}^{512}$
  - $r_i \in \mathbb{R}^{512}$



# Residuals and layer normalization

- skip/residual connection
- allow direct information flow
- (layer) normalization
- across the features
- to avoid covariate shift
- to avoid slow down training



<https://jalammar.github.io/illustrated-transformer/>

# Decoder layers

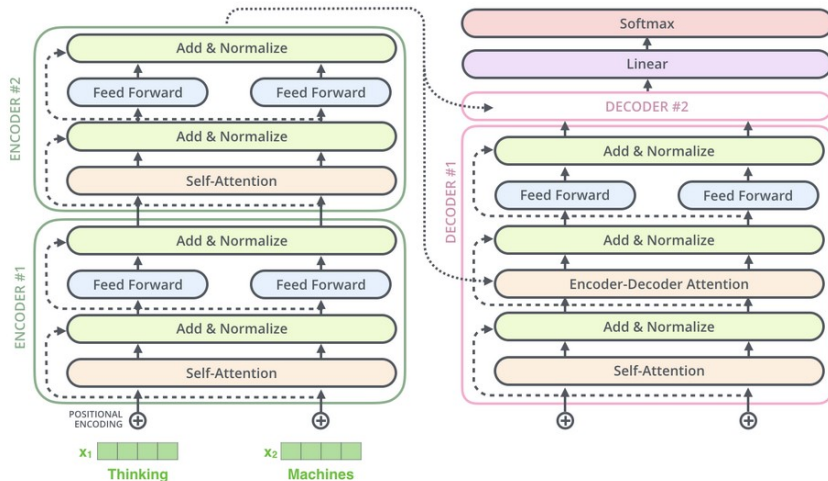
- decoder is very similar to encoder layer
- masked MHA at the bottom
- padding mask: too long or too short sequences
  - pad 0 after shorter sequences
  - pad  $\approx -\infty$  for pieces of longer sequences
- sequence mask:
  - set future positions to  $\approx -\infty$
- input to decoder is shifted right
  - prevent from just learning the copy/paste task



# Output layer

- after all decoder layers:
  - ① linear layer
  - ② softmax
- linear layer
  - fully connected NN (up-projection)
  - vocabulary size
  - results in logits
- softmax layer
  - transforms logits to pseudo-probabilities
  - then,  $\arg \max$  a typical prediction

# A transformer architecture for $N = 2$



<https://jalammar.github.io/illustrated-transformer/>

cf. also <https://medium.com/@yulemoon/detailed-explanations-of-transformer-step-by-step-dc32d90b3a98>

# Complexity comparison

Layer type	Complexity	Number of sequential operations
Attention	$O(n^2 d)$	$O(1)$
RNN	$O(nd^2)$	$O(n)$
CNN	$O(knd^2)$	$O(1)$

- $n$  is the sequence length
- $d$  is the embedding size
- $k$  is the kernel size
- often  $n < d$
- more important: Attention better parallelizable due to less sequential operations

# Step-by-step calculation of the transformer architecture<sup>1</sup> (encoder)

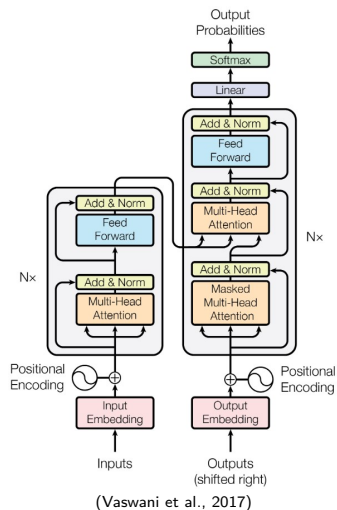
- ① data, vocabulary, vocabulary size
- ② encoding & static embedding
- ③ positional encoding
- ④ combining embedding and positional encoding
- ⑤ single-head attention (SHA) → multi-head attention (MHA)
  - ① calculating query, key, and value matrices
  - ② query and key matrix multiplication & scaling
  - ③ calculating softmax & multiplication with value matrix
  - ④ repeat, concatenate, linear transformation for final SHA/MHA output matrix
- ⑥ add & normalize
- ⑦ FFNN

<sup>1</sup>inspired by <https://levelup.gitconnected.com/>

understanding-transformers-from-start-to-end-a-step-by-step-math-example-16d4e64e6eb1

# Step-by-step calculation of the transformer architecture (decoder)

- 8 masked MHA
- 9 decoder MHA
- 10 predicting tokens



## Step 1 — data, vocabulary, vocabulary size

- assume a dataset containing (for simplicity: only) three sentences
  - ① I drink and I know things.
  - ② When you play the game of thrones, you win or you die.
  - ③ The true enemy won't wait out the storm, he brings the storm.
- tokenizing (for simplicity: without subword information) results in
  - ① I, drink, and, I, know, things
  - ② when, you, play, the, game, of, thrones, you, win, or, you, die
  - ③ the, true, enemy, won't, wait, out, the, storm, he, brings, the, storm
- the resulting set of vocabularies is given by

*I, drink, and, know, things, when, you, play, the, game, of, thrones, win, or, die, true, enemy, won't, wait, out, storm, he, brings*
- $V = 23$

## Step 2 — encoding

encoding (just mixing up the vocabularies a little bit for demonstration):

1	2	3	4	5	6	7	8	9	10	11	12
I	drink	things	know	when	won't	play	out	true	storm	brings	game
13	14	15	16	17	18	19	20	21	22	23	
the	win	of	enemy	you	wait	thrones	and	or	die	he	

we select the second sentence as an example and input the first part into the encoder of the transformer, i.e.

when	you	play	game	of	thrones
5	17	7	12	15	19

## Step 2 — static embedding

d	when 5	you 17	play 7	game 12	of 15	thrones 19
1	0.79	0.38	0.01	0.12	0.88	0.60
2	0.60	-0.37	1.93	1.73	1.24	1.02
3	0.96	0.01	0.18	0.52	0.62	0.53
4	0.64	-0.21	0.31	-0.77	-0.36	0.51
5	0.97	0.90	0.56	0.06	0.49	0.93
6	0.20	-0.26	0.59	-0.63	-0.30	0.21

- the attention paper uses  $d = 512$ , we select  $d = 6$  for demonstration
- the input embedding layer (somewhat a *lookup* table) is initialized randomly and updated during training
- the (current) static embedding for *of* is  $(0.88, 1.25, 0.04, -0.03, 0.32, -0.28)$



## Step 3 — positional encoding I

$$\text{PE}(x \mid \text{pos}) = \begin{cases} \sin(\text{pos}/10000^{x/d}) & \text{if } x \text{ is even} \\ \cos(\text{pos}/10000^{(x-1)/d}) & \text{if } x \text{ is odd} \end{cases}$$

- let us encode the position of the token *of*, which is the fifth token in our example sentence
- since programming languages usually start counting at 0, this refers to token position 4

pos	x	even/odd	formula	PE(x   pos)
4	0	even	$\sin(4/10000^{0/6})$	-0.7568
4	1	odd	$\cos(4/10000^{0/6})$	-0.6536
4	2	even	$\sin(4/10000^{2/6})$	0.1846
4	3	odd	$\cos(4/10000^{2/6})$	0.9828
4	4	even	$\sin(4/10000^{4/6})$	0.0086
4	5	odd	$\cos(4/10000^{4/6})$	1.0000

- the positional encoding for *of* is given by the PE column

## Step 3 — positional encoding II

- applying the same calculation to all input positions results in

d	when 5	you 17	play 7	game 12	of 15	thrones 19
1	0.0000	0.8415	0.9093	0.1411	-0.7568	-0.9589
2	1.0000	0.5403	-0.4161	-0.9900	-0.6536	0.2837
3	0.0000	0.0464	0.0927	0.1388	0.1846	0.2300
4	1.0000	0.9989	0.9957	0.9903	0.9828	0.9732
5	0.0000	0.0022	0.0043	0.0065	0.0086	0.0108
6	1.0000	1.0000	1.0000	1.0000	1.0000	0.9999

## Step 4 — combining embedding and positional encoding

d	when 5	you 17	play 7	game 12	of 15	thrones 19
1	0.79	0.38	0.01	0.12	0.88	0.60
2	0.60	-0.37	1.93	1.73	1.24	1.02
3	0.96	0.01	0.18	0.52	0.62	0.53
4	0.64	-0.21	0.31	-0.77	-0.36	0.51
5	0.97	0.90	0.56	0.06	0.49	0.93
6	0.20	-0.26	0.59	-0.63	-0.30	0.21

(input embedding) + (positional encoding)

=


d	when 5	you 17	play 7	game 12	of 15	thrones 19
1	0.0000	0.8415	0.9093	0.1411	-0.7568	-0.9589
2	1.0000	0.5403	-0.4161	-0.9900	-0.6536	0.2837
3	0.0000	0.0464	0.0927	0.1388	0.1846	0.2300
4	1.0000	0.9989	0.9957	0.9903	0.9828	0.9732
5	0.0000	0.0022	0.0043	0.0065	0.0086	0.0108
6	1.0000	1.0000	1.0000	1.0000	1.0000	0.9999

(input matrix for multi-head attention)


d	when 5	you 17	play 7	game 12	of 15	thrones 19
1	0.79	1.22	0.92	0.26	0.12	-0.36
2	1.60	0.17	1.51	0.74	0.59	1.30
3	0.96	0.06	0.27	0.66	0.80	0.76
4	1.64	0.79	1.31	0.22	0.62	1.48
5	0.97	0.90	0.56	0.07	0.50	0.94
6	1.20	0.74	1.59	0.37	0.70	1.21

# Step 5 — single-head attention (SHA) → multi-head attention (MHA)

Scaled Dot-Product Attention



Multi-Head Attention



- 1 calculating query, key, and value matrices
- 2 query and key matrix multiplication & scaling
- 3 calculating softmax & multiplication with value matrix
- 4 repeat, concatenate & linear transformation for final SHA/MHA output matrix

## Step 5.1 — calculating query, key, and value matrices


(input matrix for multi-head attention)							(linear weights for query $W^Q$ )				(query matrix $Q$ )			
when you play game of thrones	0.79	1.60	0.96	1.64	0.97	1.20	0.52	0.45	0.91	0.69	3.88	3.80	4.08	3.42
	1.22	0.17	0.06	0.79	0.90	0.74	0.05	0.85	0.37	0.83	2.55	1.86	2.77	1.78
	0.92	1.51	0.27	1.31	0.56	1.59	0.49	0.10	0.56	0.61	3.39	3.60	3.49	2.72
	0.26	0.74	0.66	0.22	0.07	0.37	0.71	0.64	0.40	0.14	1.02	1.18	1.24	1.30
	0.12	0.59	0.80	0.62	0.50	0.70	0.76	0.27	0.92	0.67	1.90	1.56	1.88	1.53
	-0.36	1.30	0.76	1.48	0.94	1.21	0.85	0.56	0.57	0.07	3.04	2.90	2.73	2.22
							(linear weights for key $W^K$ )				(key matrix $K$ )			
							0.74	0.57	0.21	0.73	3.71	4.04	4.15	3.41
							0.55	0.16	0.90	0.17	2.18	2.51	1.64	1.93
							0.25	0.74	0.80	0.98	3.28	3.11	3.65	3.01
							0.80	0.73	0.20	0.31	1.07	1.13	1.64	1.35
							0.37	0.96	0.42	0.08	1.49	1.97	2.14	1.81
							0.28	0.41	0.87	0.86	2.51	3.04	3.45	2.28
							(linear weights for value $W^V$ )				(value matrix $V$ )			
							0.62	0.07	0.70	0.95	3.63	4.58	4.21	4.76
							0.20	0.97	0.61	0.35	2.22	1.83	2.17	3.25
							0.57	0.80	0.61	0.50	3.12	3.77	3.41	4.33
							0.67	0.35	0.98	0.54	1.09	1.65	1.32	1.38
							0.47	0.83	0.34	0.94	1.72	2.34	1.80	2.21
							0.60	0.69	0.13	0.98	2.63	3.98	2.93	3.36

- we select  $d_k = d_v = 4$  for simplicity

## Step 5.2 — query and key matrix multiplication & scaling

(query matrix $Q$ )					(transposed key matrix $K^T$ )							$(QK^T)$					
3.88	3.80	4.08	3.42		3.71	2.18	3.28	1.07	1.49	2.51	=	58.34	31.29	49.73	19.75	28.19	43.16
2.55	1.86	2.77	1.78		4.04	2.51	3.11	1.13	1.97	3.04		34.54	18.21	29.62	11.78	16.61	25.67
3.39	3.60	3.49	2.72	$\times$	4.15	1.64	3.65	1.64	2.14	3.45		50.88	27.40	43.24	17.09	24.53	37.70
1.02	1.18	1.24	1.30		3.41	1.93	3.01	1.35	1.81	2.28		18.13	9.73	15.45	6.21	8.85	13.39
1.90	1.56	1.88	1.53									26.37	14.09	22.55	8.94	12.70	19.49
3.04	2.90	2.73	2.22									41.89	22.67	35.64	14.00	20.10	30.93

### Scaled Dot-Product Attention



- calculating  $QK^T$
- scaling with  $\sqrt{d_k} = \sqrt{4} = 2$

$(QK^T/2)$					
29.17	15.64	24.86	9.88	14.10	21.58
17.27	9.11	14.81	5.89	8.30	12.84
25.44	13.70	21.62	8.54	12.27	18.85
9.06	4.87	7.72	3.10	4.42	6.70
13.19	7.04	11.28	4.47	6.35	9.74
20.95	11.34	17.82	7.00	10.05	15.46

## Step 5.3 — calculating softmax

$$\text{softmax}(z_i) = \exp(z_i) / \sum_{j=1}^d \exp(z_j), \quad i = 1, \dots, d$$


$$\text{softmax}(29.17) = \exp(29.17) / (\exp(29.17) + \exp(15.64) + \exp(24.86) + \exp(9.88) + \exp(14.10) + \exp(21.58)) = 0.9862$$

$\text{softmax}(QK^T/2) =$	0.9862	0.0000	0.0133	0.0000	0.0000	0.0005	when
	0.9111	0.0003	0.0777	0.0000	0.0001	0.0108	you
	0.9772	0.0000	0.0214	0.0000	0.0000	0.0013	play
	0.7231	0.0108	0.1897	0.0019	0.0070	0.0676	game
	0.8450	0.0018	0.1251	0.0001	0.0009	0.0270	of
	0.9542	0.0001	0.0418	0.0000	0.0000	0.0040	thrones


## Step 5.3 — multiplication of softmax and value matrix

$\text{softmax}(QK^T/2)$						(value matrix $V$ )				$(Z)$			
0.9862	0.0000	0.0133	0.0000	0.0000	0.0005	3.63	4.58	4.21	4.76	3.6227	4.5689	4.1987	4.7536
0.9111	0.0003	0.0777	0.0000	0.0001	0.0108	2.22	1.83	2.17	3.25	3.5790	4.5095	4.1332	4.7108
0.9772	0.0000	0.0214	0.0000	0.0000	0.0013	3.12	3.77	3.41	4.33	3.6174	4.5614	4.1908	4.7485
0.7231	0.0108	0.1897	0.0019	0.0070	0.0676	1.09	1.65	1.32	1.38	3.4326	4.3353	3.9277	4.5437
0.8450	0.0018	0.1251	0.0001	0.0009	0.0270	1.72	2.34	1.80	2.21	3.5343	4.4548	4.0688	4.6626
0.9542	0.0001	0.0418	0.0000	0.0000	0.0040	2.63	3.98	2.93	3.36	3.6049	4.5439	4.1717	4.7368

Scaled Dot-Product Attention



Multi-Head Attention



- calculating  $Z = \text{softmax}(QK^T/\sqrt{d_k})V$
- this is the last step in the SHA setting
- in this example, we won't repeat this procedure  $h$  times to get MHA setting
  - if so: concatenate  $Z$  matrices



## Step 5.4 — linear transformation for final SHA/MHA output matrix

(Z)					(linear weight matrix $W^0$ )							(output matrix of multi-head attention)					
3.6227	4.5689	4.1987	4.7536	$\times$	0.80	0.34	0.45	0.54	0.07	0.53	$=$	12.33	10.87	8.07	8.71	7.12	9.16
3.5790	4.5095	4.1332	4.7108		0.85	0.74	0.78	0.50	0.75	0.55		12.18	10.73	7.97	8.60	7.02	9.05
3.6174	4.5614	4.1908	4.7485		0.53	0.81	0.55	0.59	0.49	0.14		12.32	10.85	8.06	8.70	7.10	9.14
3.4326	4.3353	3.9277	4.5437		0.70	0.60	0.12	0.42	0.29	0.87		11.69	10.28	7.63	8.25	6.73	8.71
3.5343	4.4548	4.0688	4.6626									12.03	10.59	7.86	8.49	6.93	8.95
3.6049	4.5439	4.1717	4.7368									12.27	10.81	8.03	8.67	7.08	9.11

- once again we calculate a linear transformation, here:  $ZW^0$
- dimensions of  $W^0$  have to be set in order to get an output matrix that matches the dimensions of the input

# Step 6 — add

when	0.79	1.60	0.96	1.64	0.97	1.20
you	1.22	0.17	0.06	0.79	0.90	0.74
play	0.92	1.51	0.27	1.31	0.56	1.59
game	0.26	0.74	0.66	0.22	0.07	0.37
of	0.12	0.59	0.80	0.62	0.50	0.70
thrones	-0.36	1.30	0.76	1.48	0.94	1.21

(MHA input matrix) + (MHA output matrix)

=

when	13.12	12.47	9.03	10.35	8.09	10.36
you	13.40	10.90	8.03	9.39	7.92	9.79
play	13.24	12.36	8.33	10.01	7.66	10.73
game	11.95	11.02	8.29	8.47	6.80	9.08
of	12.15	11.18	8.66	9.11	7.43	9.65
thrones	11.91	12.11	8.79	10.15	8.02	10.32

(matrix to normalize)

- the input and output matrices are just added together


# Step 6 — normalize

(matrix to normalize)								(matrix after normalization)						
when	13.12	12.47	9.03	10.35	8.09	10.36	→	when	1.3176	0.9817	-0.7957	-0.1137	-1.2814	-0.1085
you	13.40	10.90	8.03	9.39	7.92	9.79		you	1.7078	0.4862	-0.9162	-0.2516	-0.9699	-0.0562
play	13.24	12.36	8.33	10.01	7.66	10.73		play	1.3026	0.9007	-0.9402	-0.1728	-1.2463	0.1561
game	11.95	11.02	8.29	8.47	6.80	9.08		game	1.4140	0.9236	-0.5159	-0.4209	-1.3015	-0.0993
of	12.15	11.18	8.66	9.11	7.43	9.65		of	1.4270	0.8628	-0.6030	-0.3412	-1.3184	-0.0271
thrones	11.91	12.11	8.79	10.15	8.02	10.32		thrones	1.0371	1.1596	-0.8738	-0.0408	-1.3454	0.0633

- row-wise (token-wise) we determine mean  $\bar{x}$  and standard deviation  $s_x$
- we calculate a classical normalization  $\frac{x - \bar{x}}{s_x + \epsilon}$
- we here select  $\epsilon = 0.0001$

# Step 7 — FFNN

- after *add & norm*, a classical FFNN is applied to the resulting matrix
- for the example, we assume this to be very simplistic
  - here: one linear layer + ReLU activation function  $\max\{0, XW + b\}$ 
    - max is applied element-wise,  $b$  is added per row
  - realistic: multiple linear layers with activation functions



## Step 7 — FFNN linear layer

(matrix after normalization)							(weight matrix $W$ )							
when	1.3176	0.9817	-0.7957	-0.1137	-1.2814	-0.1085	$\times$	0.50	0.05	0.97	0.22	0.56	0.02	
you	1.7078	0.4862	-0.9162	-0.2516	-0.9699	-0.0562		0.17	0.52	0.63	0.48	0.06	0.60	
play	1.3026	0.9007	-0.9402	-0.1728	-1.2463	0.1561		0.53	0.87	0.47	0.10	0.31	0.79	
game	1.4140	0.9236	-0.5159	-0.4209	-1.3015	-0.0993		0.83	0.58	0.38	0.09	0.64	0.25	
of	1.4270	0.8628	-0.6030	-0.3412	-1.3184	-0.0271		0.81	0.85	0.74	0.35	0.31	0.53	
thrones	1.0371	1.1596	-0.8738	-0.0408	-1.3454	0.0633		0.25	0.31	0.22	0.77	0.57	0.85	
							(XW)							
							$=$	when	-0.7555	-1.3047	0.5073	0.1392	0.0182	-0.8130
								you	-0.5575	-1.4466	0.7066	0.1121	0.2078	-1.0226
								play	-0.8078	-1.3957	0.4355	0.2933	0.0841	-0.7473
								game	-0.8378	-1.2790	0.5661	0.1330	-0.0421	-0.7045
								of	-0.8173	-1.3315	0.5331	0.1548	0.0214	-0.7372
								thrones	-0.8552	-1.2530	0.3287	0.2716	-0.0276	-0.6433
							(XW + b)							
when	-0.7555	-1.3047	0.5073	0.1392	0.0182	-0.8130	$=$	when	-0.3355	-1.1247	0.7573	0.5592	0.3682	-0.3630
you	-0.5575	-1.4466	0.7066	0.1121	0.2078	-1.0226		you	-0.1375	-1.2666	0.9566	0.5321	0.5578	-0.5726
play	-0.8078	-1.3957	0.4355	0.2933	0.0841	-0.7473		play	-0.3878	-1.2157	0.6855	0.7133	0.4341	-0.2973
game	-0.8378	-1.2790	0.5661	0.1330	-0.0421	-0.7045		game	-0.4178	-1.0990	0.8161	0.5530	0.3079	-0.2545
of	-0.8173	-1.3315	0.5331	0.1548	0.0214	-0.7372		of	-0.3973	-1.1515	0.7831	0.5748	0.3714	-0.2872
thrones	-0.8552	-1.2530	0.3287	0.2716	-0.0276	-0.6433		thrones	-0.4352	-1.0730	0.5787	0.6916	0.3224	-0.1933
(XW) + (bias vector $b$ )														
bias	0.4200	0.1800	0.2500	0.4200	0.3500	0.4500								

## Step 7 — FFNN activation

$\max\{0, XW + b\} =$	when	0.0000	0.0000	0.7573	0.5592	0.3682	0.0000
	you	0.0000	0.0000	0.9566	0.5321	0.5578	0.0000
	play	0.0000	0.0000	0.6855	0.7133	0.4341	0.0000
	game	0.0000	0.0000	0.8161	0.5530	0.3079	0.0000
	of	0.0000	0.0000	0.7831	0.5748	0.3714	0.0000
	thrones	0.0000	0.0000	0.5787	0.6916	0.3224	0.0000

- after the FFNN there is an additional *add & norm* layer applied
- this completes the *first* encoder layer
- we won't calculate additional encoder layers here
- in practice there are  $N$  encoder layers, where the output of the *add & norm* layer (mentioned above) would serve as input for the second encoder layer

# Decoder elements

- encoder input: when you play game of thrones
- decoder input: <start> you win or you die <end>
- most of the calculation in the decoder is the same
- three elements are new:
  - ⑧ decoder MHA
  - ⑨ masked MHA
  - ⑩ predicting tokens

## Small detail: padding

- assume  $n\_sequence = 10$  as sequence length of the model
- our input sentence is of size 6
- we have to *pad* the encoder input tokens 7, 8, 9, and 10 with 0s

d	when 5	you 17	play 7	game 12	of 15	thrones 19	PAD	PAD	PAD	PAD
1	0.79	0.38	0.01	0.12	0.88	0.60	0	0	0	0
2	0.60	-0.37	1.92	1.73	1.25	1.02	0	0	0	0
3	0.96	-0.15	-0.14	0.05	0.04	-0.12	0	0	0	0
4	0.64	-0.19	0.40	-0.58	-0.03	1.01	0	0	0	0
5	0.97	0.85	0.47	-0.07	0.32	0.71	0	0	0	0
6	0.20	-0.26	0.59	-0.62	-0.28	0.24	0	0	0	0



# Padding in MHA

$$Z = \text{softmax} \left( \frac{QK^T + \text{MASK}}{\sqrt{d_k}} \right) V$$

- MASK is just a padding matrix in the encoder
- here  $-\infty$  is the padding token, since it results in 0s after softmax
- in our example:

$$\text{MASK} = \begin{array}{cccccccccc} 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \end{array}$$

## Step 8 — masked MHA

$$Z = \text{softmax} \left( \frac{QK^T + \text{MASK}}{\sqrt{d_k}} \right) V$$

- masking in the decoder MHA works quite similar to padding
- we again use  $-\infty$  as masking token
- we mask all future tokens, i.e.

$$\text{MASK} = \begin{array}{cccccccccc} 0 & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \end{array}$$


## Input for the decoder

- `<start>` and `<end>`, as well as other special tokens (e.g., `<sep>`), are part of the vocabulary set
- the decoder input matrix looks like this:

d	<code>&lt;start&gt;</code>	you	win	or	you	die	<code>&lt;end&gt;</code>	PAD	PAD	PAD
	<code>&lt;start&gt;</code>	17	14	21	17	22	<code>&lt;end&gt;</code>			
1	0.51	0.38	0.91	0.12	0.38	0.60	0.11	0	0	0
2	0.83	-0.37	1.92	0.03	-0.37	1.22	0.12	0	0	0
3	0.22	-0.15	-0.14	0.05	-0.15	-0.12	-0.50	0	0	0
4	0.04	-0.55	0.20	-0.58	-0.55	-1.01	0.01	0	0	0
5	-0.11	0.85	0.77	-0.57	0.85	0.31	1.30	0	0	0
6	0.20	-1.80	0.59	-0.62	-1.80	0.24	0.012	0	0	0

## Step 9 — decoder MHA

- $Z = \text{softmax}((QK^T + \text{MASK})/\sqrt{d_k})V$ 
  - $Q$  is calculated from the output of the first *add & norm* layer from the decoder
    - $Q = X^D W^Q$
    - where  $X^D$  is the output from the *add & norm* layer
  - $K, V$  are calculated using the output of the last encoder layer
    - $K = X^E W^K$
    - $V = X^E W^V$
    - where  $X^E$  is the output from the last encoder layer
- in the following *add & norm* layer, we calculate the normalization of  $X^D + ZW^O$



## Step 9 — decoder MHA MASK matrix

$$Z = \text{softmax} \left( \frac{QK^T + \text{MASK}}{\sqrt{d_k}} \right) V$$

- here we have a quite similar structure to encoder MHA
- $Q$  comes from decoder with one more input token  $\rightarrow$  one row less to pad with  $-\infty$

$$\text{MASK} = \begin{array}{cccccccccc} 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & 0 & 0 & 0 & 0 & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty & -\infty \end{array}$$

## Step 10 — predicting tokens

- we learn a final linear layer that projects the output of the last *add & norm* layer to logits for each token position
- output of last *add & norm* layer:  $n\_sequence \times d$  (here:  $10 \times 6$ )
- output of final linear layer:  $n\_sequence \times V$  (here:  $10 \times 23 + \text{number of special tokens}$ )
- we need a weight matrix  $W$  of size  $d \times V$

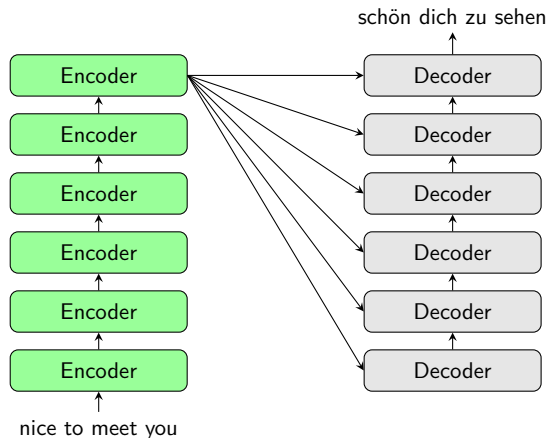
$$XW = L, \quad X \in \mathbb{R}^{n\_seq \times d}, W \in \mathbb{R}^{d \times V}, L \in \mathbb{R}^{n\_seq \times V}$$

- output of the final softmax layer is  $\text{softmax}(L) \in [0, 1]^{n\_seq \times V}$
- with  $\sum_{i=1}^V \text{softmax}(L)_{i,j} = 1 \quad \forall j = 1, \dots, n\_seq$
- these can be seen as pseudo-probabilities
- to get the encoding index for the prediction of token  $j$  simply apply  $\arg \max_{i=1, \dots, V} L_{i,j}$

## Further resources

see also <https://www.youtube.com/watch?v=EixI6t5oif0>

# Breaking BERT<sup>1</sup> down<sup>2,3</sup>



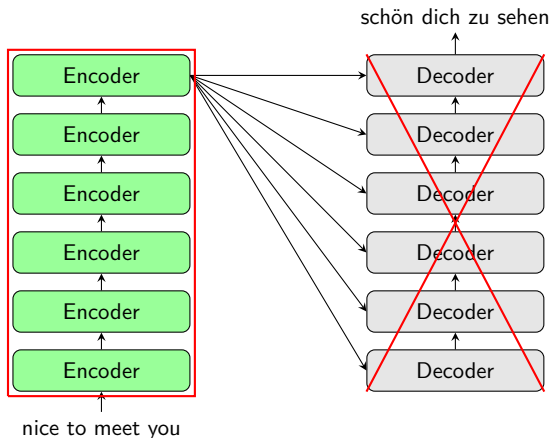
<sup>1</sup> Bidirectional Encoder Representations from Transformers

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding (Devlin et al., NAACL 2019) <http://dx.doi.org/10.18653/v1/N19-1423>

<sup>2</sup><https://jalammar.github.io/illustrated-transformer/> <sup>3</sup><https://towardsdatascience.com/breaking-bert-down-430461f60efb>



# Breaking BERT<sup>1</sup> down<sup>2,3</sup>



<sup>1</sup> Bidirectional Encoder Representations from Transformers


BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding (Devlin et al., NAACL 2019) <http://dx.doi.org/10.18653/v1/N19-1423>

<sup>2</sup><https://jalammar.github.io/illustrated-transformer/> <sup>3</sup><https://towardsdatascience.com/breaking-bert-down-430461f60efb>

# Encoder vs. decoder

- originally BERT is an **encoder only**
  - bidirectional approach contradicts traditional LM
  - predicting masked tokens (inside a sentence)
  - output: embeddings
  - extractive tasks (classification)
- there are **decoder only** models as well, e.g., GPT
  - unidirectional
  - predicting all tokens (to the right of a sentence)
  - output: (sequence of) words
  - language generation
- **encoder-decoder** models useful, e.g., for MT

# Encoder layer in BERT



- every encoder layer contains of
  - ① self-attention layer
  - ② FF layer
- input of first encoder: list of embeddings (typically  $d = 512$ )
- size of list: hyperparameter ( $\approx$  longest sequence in training data)
- input of all other encoders: output of previous encoder (same  $d$ )

# BERT's innovation

- traditional LM: left-to-right
- ELMo: left-to-right and right-to-left
- BERT: bidirectional training
- using masked language modeling (see following slides)
- also uses a subword tokenizer (WordPiece), cf. fastText
- Google released two (pre-trained) versions:
  - pre-trained encoder
  - base: 12 layers, 768 output size, 12 MHA layers, 30 522 vocab  
number of parameters: 110 million
  - large: 24 layers, 1024 output size, 16 MHA layers, 30 522 vocab  
number of parameters: 340 million

# How to train BERT

- pre-training was done on
  - BooksCorpus (800 million words)
  - English Wikipedia (2500 million words)
- took 4 days on 64 TPUs (tensor processing unit)
- TPUs are (even) faster in matrix operations than GPUs
- fine-tuning typically on single GPU

how to pre-train encoders?

# Masked language modeling (MLM)

- BERT can see all the words in the sentences
- force BERT to learn embeddings without seeing the answer
- idea:
  - ① replace a fraction of words in the input with a special [MASK] token
  - ② predict these words

# Masked language modeling (MLM)

- BERT can see all the words in the sentences
- force BERT to learn embeddings without seeing the answer
- idea:
  - ① replace a fraction of words in the input with a special [MASK] token
  - ② predict these words
- in BERT:
  - 15% of (sub)words are sampled to predict
    - 80%: replace with [MASK]
    - 10%: replace with random token
    - 10%: replace with self
  - why these values?
    - 100% [MASK] → only learning masked words' embeddings
    - 0% self → BERT would know:  
predict non-masked → always wrong word
    - 0% random: analogously to 0% self

# Next-sentence prediction


- pre-training includes a second objective
- given two sentences A and B
  - is B likely a sentence that follows A?
  - binary classification task
  - balanced, i.e., 50%/50% positive/negative examples
- idea: learn relationship between sentences
  - e.g., for question answering (QA) tasks



# Next-sentence prediction

- pre-training includes a second objective
- given two sentences A and B
  - is B likely a sentence that follows A?
  - binary classification task
  - balanced, i.e., 50%/50% positive/negative examples
- idea: learn relationship between sentences
  - e.g., for question answering (QA) tasks
- technically (cf. next slide):
  - add [CLS] token at the beginning of first sentence
  - add [SEP] at the end of each sentence
  - add segment (sentence) embedding
  - predict [CLS] token
- turns out that next-sentence prediction is not that important for training

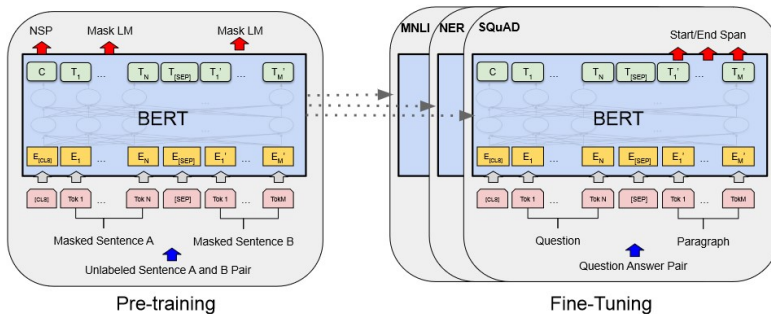
# BERT input



(Devlin et al., 2019)

- [CLS] useful for subsequent classification tasks
- [sentence A, sentence B] may be, e.g., [question, answer] for QA fine-tuning (cf. next slide)

# Fine-tuning (FT) BERT




(Devlin et al., 2019)

- several different downstream NLP tasks, e.g., NER, sentence classification, QA
- output layer for token level tasks + [CLS] classification task
- fine-tuning all parameters

# Differences in pre-training

- BERT: bidirectional
- GPT: left-to-right
- ELMo: left-to-right + right-to-left
- BERT & GPT: FT
- ELMo: feature-based



(Devlin et al., 2019)

# Transformer-based models I

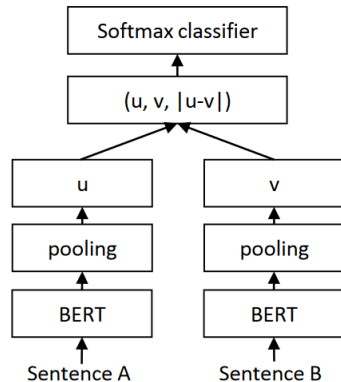
- a lot of modifications/improvements and similar architectures as BERT
- RoBERTa (English)
  - mainly just train BERT for longer and remove next sentence prediction
  - base: 12 layers, 768 output size, 50 265 vocab
  - large: 24 layers, 1024 output size, 50 265 vocab
- XLM-RoBERTa (multilingual)
  - trained on multilingual data
  - base: 12 layers, 768 output size, 250 002 vocab
  - large: 24 layers, 1024 output size, 250 002 vocab
- (m)DeBERTa (English, multilingual)
  - disentangled attention mechanism → 2 embedding vectors: content & position
  - base: 12 layers, 768 output size, 128 100 vocab (250 002 vocab)
  - large: 24 layers, 1024 output size, 128 100 vocab (250 002 vocab)

# Transformer-based models II

- DistilBERT (English)
  - less parameters, compression technique, tries to mimic BERT (base)
  - 60% parameters of BERT base, 60% faster
  - preserving over 95% of performances (GLUE benchmark)
- SBERT
  - dedicated sentence embedding training objective
  - enhanced version of predicting the [CLS] token (see following slide)
- ALBERT, CamemBERT, ELECTRA, ...

# Sentence-BERT (SBERT)

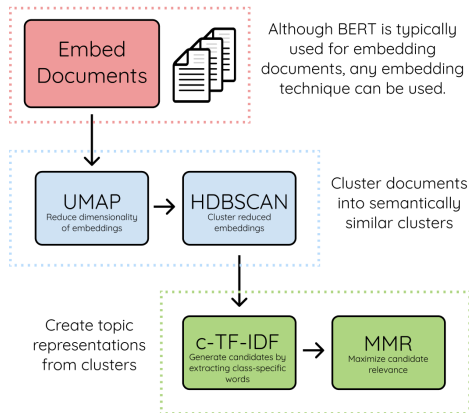
- aim: sentence embeddings
  - similar embeddings for similar sentences
- technique: siamese BERT-networks
- similarities between sentence-pairs
- classification objective:  $\text{softmax}(W_t(u, v, |u - v|))$
- triplet objective:  $\max(d(a, p) - d(a, n) + \epsilon, 0)$ 
  - $a$  anchor sentence
  - $p$  positive,  $n$  negative sentence
  - $d$  euclidean distance,  $\epsilon = 1$
- trained on STSb dataset  
[https://huggingface.co/datasets/stsb\\_multi\\_mt](https://huggingface.co/datasets/stsb_multi_mt)



<https://arxiv.org/abs/1908.10084>

Further resource: <https://www.pinecone.io/learn/sentence-embeddings/>

# How BERTopic works

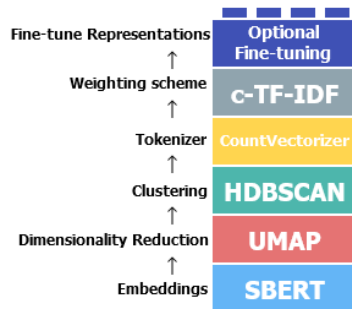


<https://github.com/MaartenGr/BERTopic/blob/master/docs/img/algorithm.png>



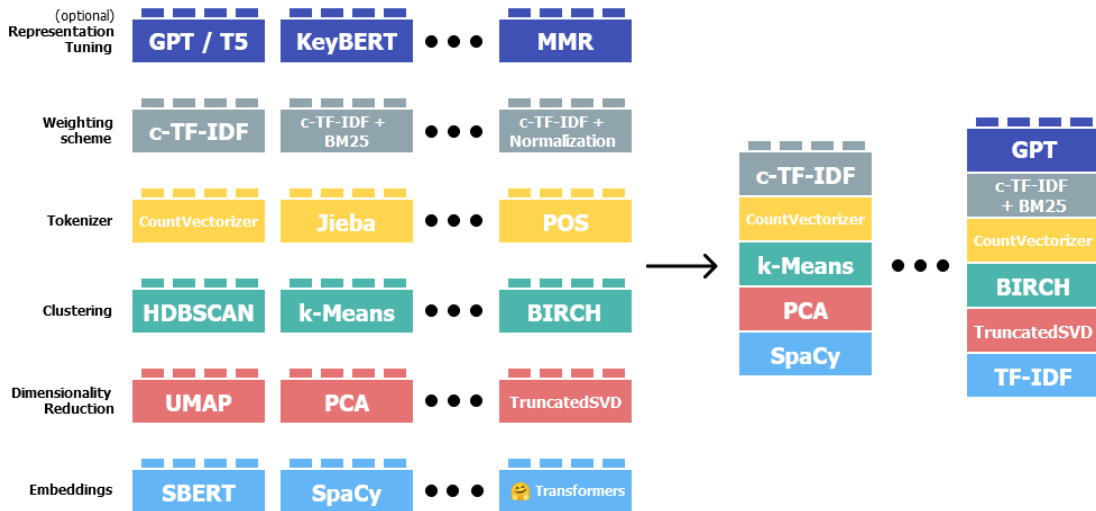
# Default workflow of BERTopic

- SBERT as embedding technique
  - <https://github.com/UKPLab/sentence-transformers>
  - default: "all-MiniLM-L6-v2"
  - non-english: "paraphrase-multilingual-MiniLM-L12-v2"
- UMAP as dimension reduction technique
  - cf., PCA, truncated SVD
- HDBSCAN as clustering technique
  - cf., k-means
- CountVectorizer for tokenization
- c-tf-idf as weighting technique
  - just a cluster-based tf-idf



<https://maartengr.github.io/BERTopic/algorithm/algorithm.html>

# Modularity of BERTopic



## c-tf-idf

$$\text{c-tf-idf}_{w,c} = \frac{\text{tf}_{w,c}}{\sum_{v=1}^V \text{tf}_{v,c}} \log \left( 1 + \frac{\frac{1}{C} \sum_{k=1}^C \sum_{v=1}^V \text{tf}_{v,k}}{\text{tf}_w} \right)$$

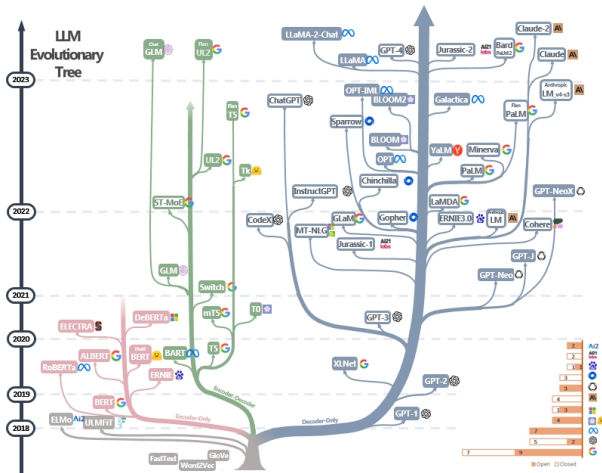
- instead of calculating idf based on documents, here based on clusters
- term-frequencies per cluster  $\text{tf}_{w,c}$  are normalized
- nominator in logarithm: average number of words per class
- actually it should be named tf-icf: term-frequency – inverse-cluster-frequency
- it is also possible to use, e.g., the importance score from slide 164

## Further resources on BERTopic

- <https://maartengr.github.io/BERTopic/>
- <https://youtu.be/uZxQz871b84>
  - Maarten Grootendorst (author) explains BERTopic
  - YouTube Video (53 min.)
- <https://youtu.be/fb7LENb9eag>
  - BERTopic explained
  - YouTube Video (45 min.)
- <https://python.plainenglish.io/topic-modeling-for-beginners-using-bertopic-and-python-aaf1b421afeb>
  - Topic Modeling For Beginners Using BERTopic and Python
  - py/medium story (member-only)
- <https://towardsdatascience.com/let-us-extract-some-topics-from-text-data-part-iv-bertopic-46ddf3c91622>
  - Let us Extract some Topics from Text Data — Part IV: BERTopic
  - tds/medium story (member-only)
- <https://towardsdatascience.com/advanced-topic-modeling-with-bertopic-85fb8a90369e>
  - Advanced Topic Modeling with BERTopic
  - tds/medium story (member-only)

# TBA

## LLM evolutionary tree



<https://github.com/Mooler0410/LLMsPracticalGuide>

- it's not all ChatGPT
- but the LLM “market” is nowadays dominated by companies
- *Part of the problem is that no university in the world today can afford to develop its own model like ChatGPT. [...] Not even us.*<sup>1</sup>
- the only remedy: stop focusing on ever larger and more powerful (and resource-hungry) models

<sup>1</sup>translated: "Teil des Problems ist schon, dass es sich keine Universität der Welt heute leisten kann, ein eigenes Modell wie ChatGPT zu entwickeln. [...] Nicht einmal wir.", Prof. Fei-Fei Li (Stanford) in *Der Spiegel* Nr. 51, 16.12.2023.

# Examples for pre-trained decoder language models

- GPT-3 (generative pretrained transformer)
  - successor of GPT and GPT-2 (12 layers, 768 output size, 50 257 vocab)
  - decoder only
  - 800GB storage for model parameters
- BLOOM (BigScience large open-science open-access multilingual language model)
  - decoder only
  - open access
  - developed as (free) alternative to GPT based models
- ChatGPT
  - at the moment (open) live study
  - no (formal) publication on it; 2023: but tech reports
  - open version an update of GPT-3 (GPT-3.5); 2023: payed access to GPT-4
- Gemini, PaLM, Claude, Falcon, LLaMA, Grok, OLMo, Llama 2, LLaMA 3.1, Mistral, Mixtral, T5, ...

## Research priorities for universities (in my opinion; not exhaustive)

What?	Why?	How?
efficiency	saving resources (time, energy, money)	parameter-efficient fine-tuning (PEFT), few-shot learning, effective learning (e.g., Adam), short-cutting via early-exit strategies
robustness	(out-of-domain/near-domain) application in real world scenarios	regularization (e.g., Adam, dropout layer), (transfer learning)
flexibility	efficient near-domain/out-of-domain application in real world scenarios	transfer learning, PEFT
reutilization	democratization of research for faster progress	transfer learning, modularity (e.g., PEFT)